

# Low level execution model

## Tools to adapt code

Whath is an execution model and how to use it ?

Henri-Pierre CHARLES

CEA DSCIN department / Grenoble

30 nov 2022

The logo for CEA (Commissariat à l'énergie atomique et aux énergies alternatives) is displayed in white lowercase letters on a red square background. Below the letters is a thick green horizontal line.



# CPU : Instruction Encoding

## Link between architecture and instruction encoding

Memory



Register Bank



Functional units



Instruction encoding



Instruction examples

```
add r1, r2, r9  
ld r3, [r3, r5]
```



Latency  
Delay

# SOA ARCH : ARM big.LITTLE

## big




Figure 2 Cortex-A15 Pipeline

## LITTLE




Figure 1 Cortex-A7 Pipeline

## How many per day ?

ARM : “more than 30 billion processors sold with more than 16M sold every day ARM” (Nov 2013)

http:

[//www.arm.com/products/processors/index.php](http://www.arm.com/products/processors/index.php)

- 4 big processors + 4 little
- Same ISA, ...
- (even for vector operation)
- Low latencie switch

big.LITTLE notion

# HWParallelismLevel uArch

## Argumentation

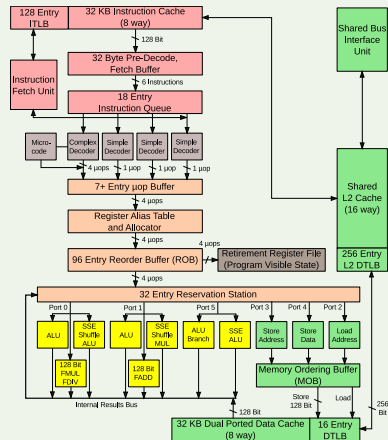
- Hidden micro architecture
- 

## Intel Example

- 700 high level instructions
- RISC internal uInstructions

## Intel Micro architecture

## Intel Core2 micro arch



Intel Core 2 Architecture

# CPU : Caches Hierarchy

- Linux command : `lstop`
- Huge impact on programs

## CPU topology

Machine (31GB total)

Package L#0

NUMANode L#0 P#0 (31GB)

L3 (8192KB)

L2 (256KB)

L2 (256KB)

L2 (256KB)

L2 (256KB)

L1d (32KB)

L1d (32KB)

L1d (32KB)

L1d (32KB)

L1i (32KB)

L1i (32KB)

L1i (32KB)

L1i (32KB)

Core L#0

Core L#1

Core L#2

Core L#3

PU L#0  
P#0


PU L#1  
P#1

PU L#2  
P#2

PU L#3  
P#3

Host: gre061041

Date: mar. 29 nov. 2022 18:05:26



# Models : C Language for Architecture


## Matrix multiply (sketch)

```
1  for (int l = 0; l < SIZE; l++)  
2    for (int c = 0; c < SIZE; c++)  
3      for (int k = 0; k < SIZE; k++)  
4        R[l][c] += A[l][k] * B[k][c];
```

## “Real world”

```
1  for (c= 0; c<NCOL; c+=cacheLineSize)  
2    for (l= 0; l<NLINE; l+=halfCacheLine)  
3      for (c2= 0; c2<NCOL; c2+=halfCacheLine)  
4        for (lk= 0; lk<halfCacheLine; lk++)  
5          for (c2k= 0; c2k<halfCacheLine; c2k++)  
6            for (ck= 0; ck<cacheLineSize; ck++)  
7              res[l+lk][c2+c2k]+= a[l+lk][c+ck] * b[c2+c2k][c+ck];
```

Learn to program = learn to serialize / schedule on defined hardware !



Other “interesting examples”

# SOA C compil chain

## Description

- Language support
  - C language (K&R), C90 up to C11
  - Normalized
  - gcc, clang, icc, tcc, ...
  - A lot of LEGACY
- Open source is mandatory : Apple, Intel
- Economic step !

## Compilation chain

- Static compilation
  - Preprocessing (# stuff)
  - Compilation
    - Lexical / Syntactic
    - IR form
    - Phases / passes
    - Instructions selection
  - Assembly
  - Load
- Dynamic part
  - OS : 1d
  - Addresses resolution
  - Cache interaction ?

[https://en.wikipedia.org/wiki/C\\_\(programming\\_language\)](https://en.wikipedia.org/wiki/C_(programming_language))



# Rappels Static Compilation chain

## Static compilation (on C language):

- |   |                    |
|---|--------------------|
| ➊ Preprocessor (all # stuff : rewriting)    | <code>cc -E</code> |
| ➋ Compilation (from C to textual assembly)  | <code>cc -S</code> |
| ➌ Assembly (from textual asm to binary asm) | <code>cc -c</code> |
| ➍ Executable (binary + dynamic library)     |                    |

## Optional

- Profiling: Compile (use `-pg`) produce File; Run File; Use `gprof`
- `cc -da` dump all intermediate representation

(Use `gcc -v` to see all the steps)

Don't stop at static time (Operating system + processor): Load in memory, dynamic linking; Branch resolution; Cache warmup

# SW-SOA GCC

## Description

- <https://gcc.gnu.org/>
- Many platforms supported
- Since 1987
- C, C++, Fortran, Objectice-C, Ada
- OpenMP, OPenACC
- GPL Licence (base compiler for Linux distro)
- Written in C, rewritten in C++ since some years

## Compilation chain

- Front-end (L to syntax tree)
- GENERIC and GIMPLE (SSA) intermediate representation
- RTL intermediate representation
- Back-end

## Features

- LTO
- Plugins
- Transactional memory support

# GCC : GCC-Research

## Code Legacy

- Invent a new architecture for new application
- Port to a new architecture
- Write C application
- Run it ... on which platform ?
- What is your metric ?

## Structure optimisation

- Invent a new optimization
  - Based on code structure
  - Based on application structure
- Use accelerator
- Performance portability


# SWParallelismLevel CUDA

- Programming language for data parallelism
- 

## Illustration

- Normalized: No
- Portable : No (NVIDIA only)
- Scalable : No, why ?

## Example of CUDA processing flow



# SW-SOA CUDA-Example

## Header and CUDA code

```
1  import pycuda.compiler as comp
2  import pycuda.driver as drv
3  import numpy
4  import pycuda.autoinit
5
6  mod = comp.SourceModule("""
7  __global__ void multiply_them(float*dest, float* a, float* b)
8  {
9      const int i=threadIdx.x;
10     dest[i]=a[i]*b[i];
11 }
12 """)
```


## Python code

```
1
2  multiply_them = mod.get_function("multiply_them")
3
4  a = numpy.random.randn(400).astype(numpy.float32)
5  b = numpy.random.randn(400).astype(numpy.float32)
6
7  dest = numpy.zeros_like(a)
8  multiply_them(
9      drv.Out(dest), drv.In(a), drv.In(b),
10     block=(400,1,1))
11
12  print dest-a*b
```

# GPU : ThreadModel

## CPU side

- Tile dataset (in cuda code)
- Map data on compute grid
- Iterate on memory move / compute



# GPU : CompilationChain

## CUDA side

- nvcc : From CUDA program to PTX (GPU assembly)
- gcc : C to binary for the host

## GPU Driver side


- CPU : Launch host process on CPU
- GPU driver : Compile PTX code to binary
- CPU : copy data from host mem to GPU mem
- CPU : map dataset and launch GPU execution
- GPU : run GPU code
- CPU : copy data from GPU mem to CPU mem

# Introduction : Remember Memory Cell 101

## SRAM memory cell depicting Inverter Loop as gates

- 6T memory cell
- Only 1 stable mode
- Read : “open” WL, read value
- Write : “open” WL, write value

## Illustration



Memory\_cell\_(computing)



# Introduction : Remember Memory 102

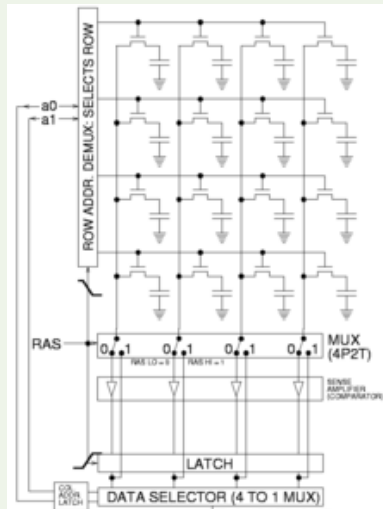
## Functions

- Select line
- Read or write
- Potentially select word in a line
- Low voltage used; "Sense amp" to normalize

Sense\_amplifier


What every programmer should know about memory

## Memory array




# Inverted Von Neumann Programming Model

## Chosen Programming model




Bus transactions

Code




```
ld r0 = @A
ld r1 = @B
add r2 = r0, r1
st r2 = @C
```



Bus transactions

Code





```
ld r0 = @ISrc1
ld r1 = @ISrc2
ld r2 = @IDst
add IDst = ISrc1, ISrc2
increment line addresses
```

## Why ?

- Allows scalability :
  - Any vector size
  - Any tile number
  - Any system configuration : near or far IMC
- Works with any processor

# Ideas : Inverted Von Neumann

## Inverted model



CPU	Memory
control flow, address compute application workload, Mem I/O	answer CPU
control flow, address compute	answer CPU (less), applica- tion workload

## Von neuman

- Stored instructions
- Bottleneck = limited bandwidth (data, insn, L1, L2, L3)
- Programm with data choregraphy


## Non von Neumann

- Break model and bottleneck
- Problems
  - Send insns
  - Synchronize
  - Data layout

# Von Neuman broke, what about Amhdal Law's ?

Ahmdal law's: "Speedup is limited by the sequential part"


## Classical approach



## Programmer approach

- Has to maximize parallel part
- Deal with data "choreography" between CPU and GPU.

## CSRAM approach



## Programmer approach

- Ease to interlace scalar instruction and IMPACT instructions
- Do not move data

# HydroGen : Initial Objectives : New Code Generation Paradigm

## Objectives: Application with Binary Automodification

- Without external libraries
- As fast as possible : Code generation speed (~10 clock cycles per instruction)
- As small as possible : Code generators (~1 KB)
- Portable across architecture : RISC-V, IBM Power 8 / 9, Kalray, CxRAM

## Benefits

- Data set code generation dependant : values, size, strides
- User programmable code generation
- Heterogeneous ISA code generation

# Domain Specific Language : HydroLang

## By The Way :

Do we need a new programming language for a new programming model ?

YES !

[Hennessy-Patterson “A New Golden Age for Computer Architecture”]

## Why a new programming language ?

- Want to make data dependent code generation, on the fly.
- Do not want to find vectorization / parallelization : use it !
- Want to **implement** “Inverted Von Neumann Model” (slides on IMC)
- Want to use specific arithmetic : integer, saturated, arithmetic, ... IP@, stochastic, geometric shapes, ...

## HydroLang features

- Similar to C syntax
- Variables = Hardware elements (register, memory line)
- Arithmetic Operators = existing processor UAL
- Run-time delayed code generation

# Motivation : Static Compiler Versus Compilette

## Static Compiler


- Run once
- Does not know data set characteristics
- Slow compilation (even with JIT)

## Compilette

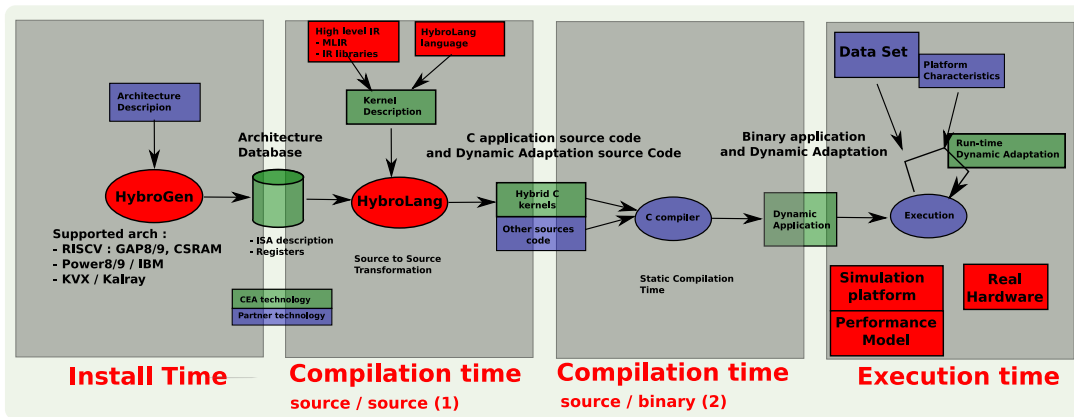
- Adapt on the fly
- Knowledge of the architecture
- Knowledge of the application

Need tools !

## Static compilation versus dynamic adaptation



# HydroGen: General View





# HydroGen : H2 Language Specific

## Language Level

- Link to hardware arithmetics datatype :int, float
- To be done : complex, ip, pixel, char
- Adapt & Link to vector size
- C like syntax
- Leaf function level
- Loop, Arithmetic expression

## Supported architecture

- RISCv : ETHZ riscy, GreenWave / GAP8/9
- CxRAM : CEA / RISCv + CXRAM
- POWER : IBM / Power8
- Kalray : Kalray / Kvx

## Run time specific

- Implement innovative code generation scenarios
- Implement innovative code optimization scenarios
- Make dataset responsive applications

# HydroGen : Objectives

## Application domains

- Stochastic number support
- Packet filtering : Datatype ipv4, ipv6 addresses
- Transprecision algorithms : usage on mathematical iterative methods
- Stencil processing


## Already done

- On the fly code generation for heterogeneous architectures
- Transprecision support : on the fly code generation for precision adaptation. Working demonstration on Newton algorithm : Power / RISCv / Kalray
- Support for In Memory Computing (next slides)
- Target processor modeling (QEMU plugin)

## Compilation

## Execution time


### (a) Static



Multiple (10) kernel executions

### (b) Dynamic

Program init



Kernel init



Application controlled



# HydroGen : Optimization Phases

## Compilation strategy

- Read code
- Gather needed instructions
- Transform in IR
- Optimize
- Generate C code generators
  - Macro instruction level
  - Instruction generators using macro instruction
  - Compilette rewriting using instruction generators

## Optimization levels

- Static :
  - Register allocation
  - Classic code optimization
- Dynamic (40 machine cycle per instructions)
  - Instruction generation
  - Instruction selection
  - Data interleaving

# HydroGen : Optimization-Target-Specific-Phases

## RISC architecture

- POWER** Specific optimization on compare & branch optimization
- Kalray** Specific optimization on arithmetic operators

## CxRAM

- Implement a code generator, generator generator
  - Detect CxRAM instructions
  - Generate RISCv instruction to generate the CxRAM instruction
  - Generate the RISCv generators

# HydroGen : Programmation

## Find a strategy

- Is your program compute bound / memory bound ?
- What parameters could be specialized ?
- What high level optimization is usefull ?
- Do you need to use a special instruction
- Regeneration frequency
  - At each function call ?
  - At each data set change ?
  - At a specific frequency ?
  - At a specific event ?

## Use Hydrogen

- Use the legacy HydroGen
- Implement a specific backend optimization
- Add a new architecture

# HydroGen : Simple Add Principle

## Not so simple example

- Code specialization / Parameter reduction
- `add(int a, int b)` by `add(int a)` for a known `b` value
- Simple, but much more efficient because at run-time

## Benefits

- Reduce memory pressure
- Reduce arithmetic (could optimize on `b` value)
- “Small” tutorial example
- Example on RISC-V (works on all architectures)

# HydroGen : Simple-Add-Source

## Simple Addition with specialization

```
1
2 typedef int (*pifi)(int);
3
4 h2_insn_t * genAdd(h2_insn_t * ptr, int b)
5 {
6     #[
7     int 32 1 add (int 32 1 a)
8     {
9         int 32 1 r;
10        r = #(b) + a; // b values will be included in code generation
11        return r;
12    }
13    ]#
14    return (h2_insn_t *) ptr;
15 }
```

# HydroGen : Simple how to Build

## How to run example

```
1 gre061041:CodeExamples/>./RunDemo.py -a riscv -i Add-With-Specialization
2 Namespace(arch=['riscv'], clean=False, debug=False, inputfile=['Add-With-Specialization'], verbose=False)
3 -->rm -f Add-With-Specialization Add-With-Specialization.c
4 -->which riscv32-unknown-elf-gcc
5 -->../HydroLang.py --toC --arch riscv --inputfile Add-With-Specialization.hl
6 -->riscv32-unknown-elf-gcc -Wall -o Add-With-Specialization Add-With-Specialization.c
7 ('3', '25')
8 -->qemu-riscv32 Add-With-Specialization 3 25
9 gre061041:CodeExamples/>qemu-riscv32 Add-With-Specialization 3 25
10 // Compilette for simple addition between 1 variable with
11 // code specialization on value = 3
12 3 + 25 = 28
```



# HydroGen : Simple-Add : Generated Source

## Code Macro instructions

```
1  #define riscv_G32(INSN){ *(h2_asm_pc++) = (INSN);}
2
3  #define RV32I_RET__I_32_1() /* RET */ \
4  do { \
5      riscv_G32(((0x8067 >> 0) & 0xffffffff)); \
6  } while(0)
7
8  #define RV32I_ADDI_RRI_I_32_1(r1,r0,i0) /* ADD */ \
9  do { \
10     riscv_G32(((i0 & 0xff) << 20)|((r0 & 0x1f) << 15)|((0x0 & 0x7) << 12)|((r1 & 0x1f) << 7)|((0x13 & 0x7f) >> 0)); \
11  } while(0)
```

# HydroGen : Simple-Add Generated Code

## Complette

```
1      h2_asm_pc = (h2_insn_t *) ptr;
2      h2_codeGenerationOK = true;
3      riscv_genMV_2(h2_00000003, (h2_sValue_t) {VALUE, 'i', 1, 32, 0, (b)}});
4      riscv_genADD_3(r, h2_00000003, a);
5      riscv_genMV_2(h2_outputVarName, r);
6      riscv_genRET_0();
7      /* Call back code for loops */
8      h2_save_asm_pc = h2_asm_pc;
9      h2_asm_pc = h2_save_asm_pc;
10     h2_iflush(ptr, h2_asm_pc);
```

# HydroGen : Simple-Add-Generated

## Instruction Selector

```
1 void riscv_genADD_3(h2_sValue_t P0, h2_sValue_t P1, h2_sValue_t P2)
2 {
3
4     if ((P0.arith == 'i') && (P0.wLen <= 32) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
5     {
6         RV32I_ADDI_RRI_I_32_1(P0.regNro, P1.regNro, P2.valueImm);
7     }
8
9     else if ((P0.arith == 'i') && (P0.wLen <= 32) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
10    {
11        RV32I_ADD_RRR_I_32_1(P0.regNro, P1.regNro, P2.regNro);
12    }
```

# HydroGen : Simple-Add Generated code

## Main code generator

```
1  h2_insn_t * genAdd(h2_insn_t * ptr, int  b)
2  {
3  /* Code Generation of 4 instructions */
4  /* Symbol table :*/
5      /*VarName = { ValOrLen, arith, vectorLen, wordLen, regNo, Value} */
6      h2_sValue_t a = {REGISTER, 'i', 1, 32, 10, 0};
7      h2_sValue_t h2_outputVarName = {REGISTER, 'i', 1, 32, 10, 0};
8      h2_sValue_t r = {REGISTER, 'i', 1, 32, 5, 0};
9      h2_sValue_t h2_00000003 = {REGISTER, 'i', 1, 32, 6, 0};
10
11
12  /* Label table :*/
13  #define riscv_genLABEL(LABEL_ID) labelAddresses[LABEL_ID] = h2_asm_pc;
14  h2_insn_t * labelAddresses []={
15      };
16
17
18      h2_asm_pc = (h2_insn_t *) ptr;
```

# HydroGen : Simple-Add-Exec

## Simple run

```
1 gre061041:CodeExamples/>qemu-riscv32 Add-With-Specialization 3 40
2 // Compilette for simple addition between 1 variable with
3 // code specialization on value = 3
4 3 + 40 = 43
```

## Run with debug

```
1 qemu-riscv32 Add-With-Specialization 3 25
2 // Compilette for simple addition between 1 variable with
3 // code specialization on value = 3
4 0x19008 : RV32I_MV_RI_I_32_1
5 0x1900c : RV32I_ADD_RRR_I_32_1
6 0x19010 : RV32I_MV_RR_I_32_1
7 0x19014 : RV32I_RET_I_32_1
8 3 + 25 = 28
```

# HydroGen : Simple-Add-Source

## Simple Addition with specialization

```
1  int main(int argc, char * argv[])
2  {
3      h2_insn_t * ptr;
4      int in0, in1, res;
5      pifi fPtr;
6
7      if (argc < 3)
8      {
9          printf("Give 2 values\n");
10         exit(-1);
11     }
12     in0 = atoi (argv[1]);    // Get the users values in1 & in2
13     in1 = atoi (argv[2]);
14     ptr = h2_malloc (1024);  // Allocate memory for 1024 instructions
15     printf("// Compilette for simple addition between 1 variable with\n");
16     printf("// code specialization on value=%d\n", in0);
17     fPtr = (pifi) genAdd (ptr, in0); // Generate instructions
18     res = fPtr(in1); // Call generated code
19     printf("%d+%d=%d\n", in0, in1, res);
```

# HydroGen : Simple-Add Debug

## 1 shell to Run / interact

```
1 gre061041:CodeExamples/>qemu-riscv32 -g \  
2 7777 Add-With-Specialization 3 25  
3 // Compilette for simple addition  
4 // between 1 variable with  
5 // code specialization on value = 3  
6 0x19008 : RV32I_MV_RI_I_32_1  
7 0x1900c : RV32I_ADD_RRR_I_32_1  
8 0x19010 : RV32I_MV_RR_I_32_1  
9 0x19014 : RV32I_RET__I_32_1
```

## 1 shell to Debug / observe

```
1 riscv32-unknown-elf-gdb Add-With-Specialization  
2 GNU gdb (GDB) 9.2  
3 ../..  
4 (gdb) target remote :7777  
5 (gdb) break main  
6 Breakpoint 1 at 0x107c6: file Add-With-Specialization.c, line 201.  
7 (gdb) c  
8 Continuing.  
9 Breakpoint 1, main (argc=3, argv=0x40800374) at Add-With-Specialization.c:201  
10 201 if (argc < 3)  
11 (gdb) n  
12 206 in0 = atoi (argv[1]);  
13 // Get the users values in1 & in2  
14 211 fPtr = (pifi) genAdd (ptr, in0); // Generate instructions  
15 (gdb)  
16 212 res = fPtr(in1); // Call generated code  
17 (gdb) x/4i fPtr  
18 0x19008: ori t1,zero,3  
19 0x1900c: add t0,t1,a0  
20 0x19010: mv a0,t0  
21 0x19014: ret  
(gdb)
```

# HydroGen : Transprecision Source H2

## Transprecision square root source code

```
1  /* Newton square root demonstration with variable precision */
2  h2_insn_t * genIterate(h2_insn_t * ptr, int FloatWidth)
3  {
4      #[
5          flt #(FloatWidth) 1 iterate(flt #(FloatWidth) 1 u, flt #(FloatWidth) 1 val, flt #(FloatWidth) 1 div )
6          {
7              flt #(FloatWidth) 1 r, tmp1, tmp2;
8              tmp1 = val / u;
9              tmp2 = u + tmp1;
10             return tmp2 / div;
11         }
12     ]#      /* r = (u + (#(value) / u)) / 2.0*/
13
14     return (h2_insn_t *) ptr;
15 }
```



# HydroGen : Transprecision Generated Macros

## Macro instruction generation

```
1  /* Begin Header autogenerated part */
2  #include "h2-power-power.h"
3  #define power_G32(INSN){ *(h2_asm_pc++) = (INSN);}
4
5  void P1_BLR__I_32(void){ /* ret */
6  #ifdef H2_DEBUG
7      printf("%p: P1_BLR__I_32\n", h2_asm_pc);
8  #endif
9      power_G32(((0x4e800020 >> 0) & 0xffffffff)); \
10 }
11
12 void PPC_FADDS_RRR_F_32(int r0, int r1, int r2){ /* add */
13 #ifdef H2_DEBUG
14     printf("%p: PPC_FADDS_RRR_F_32\n", h2_asm_pc);
15 #endif
16     power_G32(((0x3b & 0x3f) << 26)|((r0 & 0x1f) << 21)|((r1 & 0x1f) << 16)|((r2 & 0x1f) << 11)|((0x2a & 0x7ff) >> 0)); \
17 }
```

# HybroGen : Transprecision-Generated-Complette

## Complette Generation

```
1  /* Newton square root demonstration with variable precision */
2  h2_insn_t * genIterate(h2_insn_t * ptr, int FloatWidth)
3  {
4  /* Code Generation of 4 instructions */
5  /* Symbol table :*/
6      /*VarName = { ValOrLen, arith, vectorLen, wordLen, regNo, Value} */
7      h2_sValue_t u = {REGISTER, 'f', 1, (FloatWidth), 1, 0};
8      h2_sValue_t val = {REGISTER, 'f', 1, (FloatWidth), 2, 0};
9      h2_sValue_t div = {REGISTER, 'f', 1, (FloatWidth), 3, 0};
10     h2_sValue_t h2_outputVarName = {REGISTER, 'f', 1, (FloatWidth), 1, 0};
11     h2_sValue_t r = {REGISTER, 'f', 1, (FloatWidth), 14, 0};
12     h2_sValue_t tmp1 = {REGISTER, 'f', 1, (FloatWidth), 15, 0};
13     h2_sValue_t tmp2 = {REGISTER, 'f', 1, (FloatWidth), 16, 0};
14     h2_sValue_t h2_00000000 = {REGISTER, 'f', 1, (FloatWidth), 17, 0};
```

# HybroGen : Transprecision-Generated-Complette

## Complette Generation Code generator

```
1      h2_asm_pc = (h2_insn_t *) ptr;
2      h2_codeGenerationOK = 1;
3      power_genDIV_3(h2_00000000, val, u);
4      power_genMV_2(tmp1, h2_00000000);
5      power_genADD_3(h2_00000000, u, tmp1);
6      power_genMV_2(tmp2, h2_00000000);
7      power_genDIV_3(h2_00000000, tmp2, div);
8      power_genMV_2(h2_outputVarName, h2_00000000);
9      power_genRET_0();
10     /* Call back code for loops */
11     h2_save_asm_pc = h2_asm_pc;
12     h2_asm_pc = h2_save_asm_pc;
13     iflush(ptr, h2_asm_pc);
14     /* r = (u + (#(value) / u)) / 2.0*/
```

# HydroGen : Transprecision-Generated-InstructionSelector

## Macro instruction generation

```
1 void power_genADD_3(h2_sValue_t P0, h2_sValue_t P1, h2_sValue_t P2)
2 {
3     if ((P0.arith == 'f') && (P0.wLen <= 32) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
4     {
5         PPC_FADDS_RRR_F_32(P0.regNro, P1.regNro, P2.regNro);
6     }
7     else if ((P0.arith == 'f') && (P0.wLen <= 32) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
8     {
9         PPC_FADDS_RRR_F_32(P0.regNro, P1.regNro, P2.regNro);
10    }
11    else if ((P0.arith == 'f') && (P0.wLen <= 32) && (P0.vLen == 4) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
12    {
13        V2_03_VADDFP_RRR_F_32(P0.regNro, P1.regNro, P2.regNro);
14    }
15    else if ((P0.arith == 'f') && (P0.wLen <= 64) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
16    {
17        P1_FADD_RRR_F_64(P0.regNro, P1.regNro, P2.regNro);
18    }
19    else if ((P0.arith == 'f') && (P0.wLen <= 64) && (P0.vLen == 1) && P0.ValOrReg == REGISTER && P1.ValOrReg == REGISTER && P2.ValOrReg == REGISTER)
20    {
21        P1_FADD_RRR_F_64(P0.regNro, P1.regNro, P2.regNro);
22    }
23
24    else
```

# HydroGen : Transprecision-Exec

## Macro instruction generation

```
1  qemu-ppc64le Newton.power 65536 1e-13
2  Compute square root of 65536.000000
3  With precision of 1.000000e+01 (float)
4  With precision of 1.000000e-13 (double)
5  0x100212a0 : PPC_FDIVS_RRR_F_32
6  0x100212a4 : P1_FMR_RR_F_32
7  0x100212a8 : PPC_FADDS_RRR_F_32
8  0x100212ac : P1_FMR_RR_F_32
9  0x100212b0 : PPC_FDIVS_RRR_F_32
10 0x100212b4 : P1_FMR_RR_F_32
11 0x100212b8 : P1_BLR_I_32
12   0 float : 32768.500000000000000000000000, 1.000000e+01
13   1 float : 16385.250000000000000000000000, 1.000000e+01
14   2 float : 8194.625000000000000000000000, 1.000000e+01
15   3 float : 4101.3110351562500000000000, 1.000000e+01
16   4 float : 2058.6452636718750000000000, 1.000000e+01
17   5 float : 1045.2398681640625000000000, 1.000000e+01
18   6 float : 553.9696655273437500000000, 1.000000e+01
```

# HydroGen : Transprecision Exec 2nd part

## Macro instruction generation

```
1      7 float   : 336.13607788085937500000, 1.000000e+01
2      8 float   : 265.55236816406250000000, 1.000000e+01
3      9 float   : 256.17181396484375000000, 1.000000e+01
4      0x100212a0 : P1_FDIV_RRR_F_64
5      0x100212a4 : P1_FMR_RR_F_64
6      0x100212a8 : P1_FADD_RRR_F_64
7      0x100212ac : P1_FMR_RR_F_64
8      0x100212b0 : P1_FDIV_RRR_F_64
9      0x100212b4 : P1_FMR_RR_F_64
10     0x100212b8 : P1_BLR_I_32
11     10 double  : 256.00005761765521583584, 1.000000e-13
12     11 double  : 256.00000000000648014975, 1.000000e-13
13     12 double  : 256.00000000000000000000, 1.000000e-13
14     13 double  : 256.00000000000000000000, 1.000000e-13
```

# HydroGen : Transprecision Source Main

## Transprecision square root source code (main control)

```
1  fPtr1 = (piff) genIterate (ptr, FLOAT);
2  do
3  {
4      if ((diff < precf) && isFloat)
5      {
6          /* Code re-generation with double for better precision */
7          fPtr2 = (pidd) genIterate (ptr, DOUBLE);
8          isFloat = False;
9      }
10     value = next;
11     next = (isFloat)?fPtr1(value, af, 2.0):fPtr2(value, af, 2.0);
12     diff = ABS(next - value);
```

## Conclusion :

### Have fun with binary code generation

- Many scenario to be invented

### Interesting period

- End of the Moore's Law
- Amdahl law still valid
- Energy challenge = Specialized architecture
- Specialized architecture = Fun with compilers